

ECE684: NLP Final Project

Harsh Bandhey

In [166...

```
import nltk
import random
import pickle
import numpy as np
import pandas as pd
from spacy import displacy
from tqdm.notebook import tqdm
from tqdm.keras import TqdmCallback
from sklearn.metrics import classification_report
from keras.utils import pad_sequences
from keras.models import Sequential
from keras.utils.vis_utils import plot_model
from keras.layers import Dense, Dropout, TimeDistributed
from keras.layers import Embedding, LSTM, Bidirectional, CuDNNLSTM
from keras.callbacks import TensorBoard
from nltk.probability import LidstoneProbDist
from nltk.probability import SimpleGoodTuringProbDist, WittenBellProbDist
import matplotlib.pyplot as plt
import seaborn as sns
sns.set_theme()
```

Helper Functions

In [167...

```
##@title Data Loading Functions
def conll_sentences(conll_file):
    sent = []
    pos = []
    chunk = []
    entity = []
    temp_sent = []
    temp_pos = []
    temp_chunk = []
    temp_entity = []

    with open(conll_file) as f:
        conll_raw_data = f.readlines()
    conll_raw_data = [x.strip() for x in conll_raw_data]

    for line in conll_raw_data:
        if line != '':
            split_line = line.split()
            if len(split_line) == 4:
                if split_line[0] != '-DOCSTART-':
                    temp_sent.append(split_line[0])
                    temp_pos.append(split_line[1])
                    temp_chunk.append(split_line[2])

    # Rename entity values as PER, LOC, ORG, MISC, O
```

```

        old_ent = split_line[3]
        if old_ent in ('I-ORG', 'B-ORG'):
            new_ent = 'ORG'
        elif old_ent in ('I-LOC', 'B-LOC'):
            new_ent = 'LOC'
        elif old_ent in ('I-MISC', 'B-MISC'):
            new_ent = 'MISC'
        elif old_ent in ('I-PER', 'B-PER'):
            new_ent = 'PER'
        else:
            new_ent = 'O'
        temp_entity.append(new_ent)
    else:
        raise IndexError('Line split length does not equal 4.')
else:
    if len(temp_sent) > 0:
        assert(len(sent) == len(pos))
        assert(len(sent) == len(chunk))
        assert(len(sent) == len(entity))
        sent.append(temp_sent)
        pos.append(temp_pos)
        chunk.append(temp_chunk)
        entity.append(temp_entity)
        temp_sent = []
        temp_pos = []
        temp_chunk = []
        temp_entity = []

return sent, pos, chunk, entity

```

In [168...

```

#@title Evaluation Function
def get_preds(model, test_data):
    preds_list, actual_list = list(), list()
    for X in tqdm(test_data):
        sent = [i[0] for i in X]
        tagged = model.tag(sent)
        preds = [i[1] for i in tagged]
        act = [i[1] for i in X]
        preds_list += preds
        actual_list += act
    return preds_list, actual_list

def accuracy(expected, predicted):
    total = 0
    correct = 0
    for i in range(len(expected)):
        total += 1
        if (expected[i] == predicted[i]):
            correct += 1
    acc = correct/total
    print('accuracy = {} / {} = {}'.format(correct, total, round(acc*100,2)))
    return acc

def dispac_doc(item, title):
    cur = 0
    ents = []
    for a, b in item:
        ents.append({"start":cur,
                     "end":cur+len(a),

```

```

        "label": b })
    cur = cur+len(a)+1
    element = { "text": " ".join([a for a,b in item]),
                "ents": ents,
                "title": title }
    # displacy.render(doc, style="ent", manual=True, jupyter=True)
    return element

```

In [169...

```

#@title Encoding
def encode(entities_list):
    encoded_list = []
    # one-hot formatting: [PER LOC ORG MISC O]
    for entities in entities_list:
        encoded_vectors = []
        for ent in entities:
            if ent == 'PER':
                encoded_vectors.append([1, 0, 0, 0, 0])
            elif ent == 'LOC':
                encoded_vectors.append([0, 1, 0, 0, 0])
            elif ent == 'ORG':
                encoded_vectors.append([0, 0, 1, 0, 0])
            elif ent == 'MISC':
                encoded_vectors.append([0, 0, 0, 1, 0])
            else:
                encoded_vectors.append([0, 0, 0, 0, 1])
        encoded_list.append(encoded_vectors)
    return encoded_list

def decode(scores_list):
    predictions_list = []
    for predictions in scores_list:
        decoded = []
        for pred in predictions:
            max_index = np.argmax(pred)
            if max_index == 0:
                decoded.append('PER')
            elif max_index == 1:
                decoded.append('LOC')
            elif max_index == 2:
                decoded.append('ORG')
            elif max_index == 3:
                decoded.append('MISC')
            elif max_index == 4:
                decoded.append('O')
        predictions_list.append(decoded)
    return predictions_list

```

In [170...

```

#@title Glove Emedding Functions
def load_glove_dict(glove_file):
    word_dict = {}
    with open(glove_file, 'r') as f:
        for line in f:
            split = line.split()
            word = split[0]
            vector = np.array([float(v) for v in split[1:]])
            word_dict[word] = vector
    return word_dict

```

```
def get_glove_vector(g_dict, word):
    try:
        vector = g_dict[word.lower()]
    except KeyError:
        vector_len = len(g_dict['test'])
        vector = np.array([0.]*vector_len)
    return vector
```

In [171]:

```
##@title Plot Functions
def plot_acc(history):
    plt.plot(history.history['accuracy'])
    plt.plot(history.history['val_accuracy'])
    plt.title('Bi-LSTM Accuracy')
    plt.ylabel('Accuracy')
    plt.xlabel('Epochs')
    plt.legend(['train', 'Test'], loc='upper left')
    plt.show()

def plot_loss(history):
    plt.plot(history.history['loss'])
    plt.plot(history.history['val_loss'])
    plt.title('Bi-LSTM Loss')
    plt.ylabel('Loss')
    plt.xlabel('Epochs')
    plt.legend(['Train', 'Test'], loc='upper right')
    plt.show()
```

Data Processing

In [7]:

```
!mkdir dataset
!wget -q "https://drive.google.com/uc?export=download&id=1FQ8ZBCYxxrhGOhnniyi3Qzo"
!wget -q "https://drive.google.com/uc?export=download&id=1lf3hKR4ndPbJcw3QmMI-Uw"
!wget -q "https://drive.google.com/uc?export=download&id=1EhIJKhIA4DaBTcTJpszPGZ"
!wget -q "https://drive.google.com/uc?export=download&id=1mOBEvOino44PZ8UR5d0lR1"
```

In [6]:

```
train_file = './dataset/eng.train'
testa_file = './dataset/eng.testa'
testb_file = './dataset/eng.testb'
testc_file = './dataset/eng.testc'
```

In [7]:

```
entity_set = {'PER', 'LOC', 'ORG', 'MISC', 'O'}
train_sent, train_pos, train_chunks, train_entities = conll_sentences(train_file)
testa_sent, testa_pos, testa_chunks, testa_entities = conll_sentences(testa_file)
testb_sent, testb_pos, testb_chunks, testb_entities = conll_sentences(testb_file)
testc_sent, testc_pos, testc_chunks, testc_entities = conll_sentences(testc_file)
```

In [8]:

```
combined_sentences = train_sent + testa_sent + testb_sent + testc_sent
word_set = set()
for sent in combined_sentences:
    for word in sent:
        word_set.add(word)
```

```
In [9]: train_data = list(zip(train_sent, train_entities))
train_data = [list(zip(x,y)) for x,y in train_data]
```

```
In [10]: test_sent, test_entities = testa_sent+testb_sent, testa_entities+testb_entities
test_data = list(zip(test_sent, test_entities))
test_data = [list(zip(x,y)) for x,y in test_data]
```

HMM Model

Training and Evaluation

```
In [67]: trainer = nltk.tag.HiddenMarkovModelTrainer(states=entity_set, symbols=word_set)
estimator = lambda fd, bins: LidstoneProbDist(fd, 0.1, bins)
model = trainer.train_supervised(train_data, estimator=estimator)
```

```
In [109... y_true, y_pred = get_preds(model, test_data)
print(classification_report(y_true, y_pred))
```

	precision	recall	f1-score	support
LOC	0.80	0.88	0.84	3665
MISC	0.75	0.86	0.80	1903
O	0.99	0.95	0.97	84102
ORG	0.67	0.85	0.74	3611
PER	0.68	0.89	0.77	4516
accuracy			0.94	97797
macro avg	0.78	0.89	0.83	97797
weighted avg	0.95	0.94	0.95	97797

Running a few examples on the HMM Model

```
In [78]: item = test_data[3]
doc = dispacey_doc(item, "Original NER")
html = displacy.render(doc, style="ent", manual=True, jupyter=True)
```

Original NER

Their o stay o on o top o , o though o , o may o be o
 short-lived o as o title o rivals o Essex ORG , o Derbyshire ORG
 and o Surrey ORG all o closed o in o on o victory o while o
 Kent ORG made o up o for o lost o time o in o their o
 rain-affected o match o against o Nottinghamshire ORG . o

In [79]:

```
_y_pred = model.tag([a for a,b in item])
doc = displacy_doc(_y_pred, "Predicted NER")
html = displacy.render(doc, style="ent", manual=True, jupyter=True)
```

Predicted NER

Their o stay o on o top o , o though o , o may o be o
short-lived o as o title o rivals o Essex **ORG** , o Derbyshire **ORG**
and **ORG** Surrey **ORG** all o closed o in o on o victory o while
o Kent **ORG** made o up o for o lost o time o in o their o
rain-affected o match o against o Nottinghamshire **ORG** . o

Data Generation From HMM

In [68]:

```
synth_train_data = [model.random_sample(random.Random(),
                                         np.random.randint(10, 25)) for i in tqdm(range(len(train_data)))]
synth_test_data = [model.random_sample(random.Random(),
                                         np.random.randint(10, 25)) for i in tqdm(range(len(test_data)))]
```

In [76]:

```
with open('syndata.pkl', 'wb') as f:
    obj_dict = {"synth_train_data":synth_train_data,
               "synth_test_data":synth_test_data}
    pickle.dump(obj_dict, f)
```

In [70]:

```
with open('syndata.pkl', 'rb') as f:
    obj_dict = pickle.load(f)
synth_train_data = obj_dict["synth_train_data"]
synth_test_data = obj_dict["synth_test_data"]
```

In [28]:

```
item = synth_train_data[2]
doc = displacy_doc(item, "Generated NER Data")
html = displacy.render(doc, style="ent", manual=True, jupyter=True)
```

Generated NER Data

Commonwealth **ORG** 17 o with o either o billions o . o were o
military o who o The o companies o

Testing HMM on Synthetic Data

In [53]:

```
trainer = nltk.tag.HiddenMarkovModelTrainer(states=entity_set, symbols=word_set)
estimator = lambda fd, bins: SimpleGoodTuringProbDist(fd, bins=1e5)
model = trainer.train_supervised(synth_train_data, estimator=estimator)
```

```
/usr/local/lib/python3.8/dist-packages/nltk/probability.py:1452: UserWarning: SimpleGoodTuring did not find a proper best fit line for smoothing probabilities of occurrences. The probability estimates are likely to be unreliable.
warnings.warn(
```

In [102...

```
y_true, y_pred = get_preds(model, synth_test_data)
print(classification_report(y_true, y_pred))
```

	precision	recall	f1-score	support
LOC	0.46	0.75	0.57	265
MISC	0.39	0.57	0.46	197
O	0.97	0.93	0.95	9866
ORG	0.55	0.64	0.59	466
PER	0.58	0.72	0.64	497
accuracy			0.90	11291
macro avg	0.59	0.72	0.64	11291
weighted avg	0.92	0.90	0.90	11291

In [165...

```
idx = 3
item = synth_test_data[idx]
doc = displacy_doc(item, "Original NER")
html = displacy.render(doc, style="ent", manual=True, jupyter=True)
```

Original NER

finally o the o Two o of o , o (o Friday o AND o is o
who o the o Jerry PER Bevan PER David PER 0.38 o Hungary
LOC of o the o surprise o

In [164...

```
item = synth_test_data[idx]
_y_pred = model.tag([a for a,b in item])
doc = displacy_doc(_y_pred, "Predicted NER")
html = displacy.render(doc, style="ent", manual=True, jupyter=True)
```

Predicted NER

finally o the o Two o of o , o (o Friday o AND o is o
who o the o Jerry PER Bevan PER David PER 0.38 o Hungary
LOC of o the o surprise o

Netural Network (BiLSTM) Model

Data Prep for LSTM

```
In [172...
max_sent_len = len(max(combined_sentences, key=len))
word_dict = dict()
for i, wrd in enumerate(word_set):
    word_dict[wrd] = int(i) + 1
```

```
In [173...
max_features = len(word_set) + 1
maxlen = max_sent_len
batch_size = 256
output_dim = 100
epochs = 30
drop_out = 0.05
n_tags = 5
```

```
In [174...
def word2vec(sentences):
    out = []
    for sent in sentences:
        temp = []
        for word in sent:
            temp.append(word_dict[word])
        out.append(temp)
    return np.array(out)
```

```
In [175...
test_sent = testa_sent+testb_sent
train_sent_vec = pad_sequences(word2vec(train_sent), maxlen=maxlen, padding='pos
test_sent_vec = pad_sequences(word2vec(test_sent), maxlen=maxlen, padding='post'
test_entities = testa_entities+testb_entities
train_ent_vec = pad_sequences(encode(train_entities), maxlen=maxlen, padding='po
test_ent_vec = pad_sequences(encode(test_entities), maxlen=maxlen, padding='post
```

<ipython-input-174-74d0e4c39a73>:8: VisibleDeprecationWarning: Creating an ndarray from ragged nested sequences (which is a list-or-tuple of lists-or-tuples-or ndarrays with different lengths or shapes) is deprecated. If you meant to do this, you must specify 'dtype=object' when creating the ndarray.

```
return np.array(out)
```

LSTM Models

```
In [20]: !wget https://worksheets.codalab.org/rest/bundles/0x4090ba96b8a444c2a44b2c47884c
```

```
--2022-12-18 13:33:44-- https://worksheets.codalab.org/rest/bundles/0x4090ba96b8a444c2a44b2c47884c25f2/contents/blob/glove.twitter.27B.100d.txt
Resolving worksheets.codalab.org (worksheets.codalab.org)... 13.68.212.115
Connecting to worksheets.codalab.org (worksheets.codalab.org)|13.68.212.115|:443... connected.
HTTP request sent, awaiting response... 200 OK
Syntax error in Set-Cookie: codalab_session=""; expires=Thu, 01 Jan 1970 00:00:00 GMT; Max-Age=-1; Path=/ at position 70.
Length: unspecified [text/plain]
Saving to: 'glove.twitter.27B.100d.txt.1'
```

```
glove.twitter.27B.1      [          <=>          ] 974.34M   104MB/s   in 9.7s
```

```
2022-12-18 13:33:54 (100 MB/s) - 'glove.twitter.27B.100d.txt.1' saved [102166937
```


In [33]:

```
def get_glove_weights():
    glove_file='glove.twitter.27B.100d.txt'
    g_dict = load_glove_dict(glove_file)
    embedding_matrix = np.zeros((max_features, output_dim))
    for word, i in tqdm(word_dict.items()):
        embedding_vector = g_dict.get(word)
        if embedding_vector is not None:
            # words not found in embedding index will be all-zeros.
            embedding_matrix[i] = embedding_vector
    return embedding_matrix
```

In [34]:

```
embedding_matrix = get_glove_weights()
```

In [35]:

```
def get_bilstm_lstm_model():
    model = Sequential()
    model.add(Embedding(max_features, output_dim, input_length=maxlen,
                        weights=[embedding_matrix], mask_zero=True,
                        trainable=False))
    model.add(Bidirectional(LSTM(units=output_dim,
                                return_sequences=True, dropout=drop_out,
                                recurrent_dropout=drop_out),
                            merge_mode = 'concat'))
    model.add(LSTM(units=output_dim, return_sequences=True, dropout=drop_out*2,
                    recurrent_dropout=drop_out*2))
    model.add(TimeDistributed(Dense(n_tags, activation="softmax")))
    model.compile(loss='categorical_crossentropy', optimizer='adam', metrics=['a
    # model.summary()
    return model
```

Model Training

In [36]:

```
GLOVE = True
batch_size = 512
```

In [37]:

```
if GLOVE:
    model = get_bilstm_lstm_model()
    train_vec = train_sent_vec
    test_vec = test_sent_vec
else:
    pass
```

WARNING:tensorflow:Layer lstm will not use cuDNN kernels since it doesn't meet the criteria. It will use a generic GPU kernel as fallback when running on GPU.
 WARNING:tensorflow:Layer lstm will not use cuDNN kernels since it doesn't meet the criteria. It will use a generic GPU kernel as fallback when running on GPU.
 WARNING:tensorflow:Layer lstm will not use cuDNN kernels since it doesn't meet the criteria. It will use a generic GPU kernel as fallback when running on GPU.
 WARNING:tensorflow:Layer lstm_1 will not use cuDNN kernels since it doesn't meet the criteria. It will use a generic GPU kernel as fallback when running on GPU.

In [38]:

```
model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
embedding (Embedding)	(None, 124, 100)	3029000
bidirectional (Bidirectional)	(None, 124, 200)	160800
lstm_1 (LSTM)	(None, 124, 100)	120400
time_distributed (TimeDistributed)	(None, 124, 5)	505

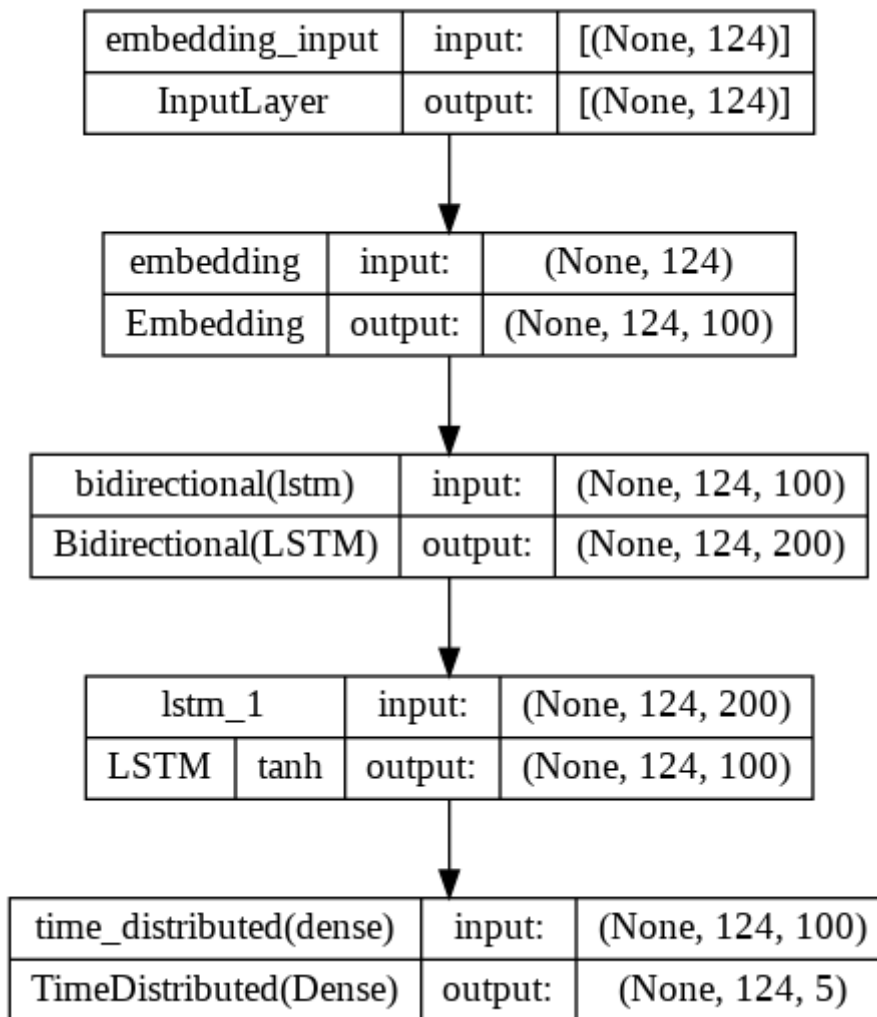
=====
Total params: 3,310,705
Trainable params: 281,705
Non-trainable params: 3,029,000
=====

In [181]:

```
plot_model(model,  
            show_shapes=True,  
            show_layer_activations=True)
```

WARNING:tensorflow:Layer lstm will not use cuDNN kernels since it doesn't meet the criteria. It will use a generic GPU kernel as fallback when running on GPU.
WARNING:tensorflow:Layer lstm will not use cuDNN kernels since it doesn't meet the criteria. It will use a generic GPU kernel as fallback when running on GPU.
WARNING:tensorflow:Layer lstm will not use cuDNN kernels since it doesn't meet the criteria. It will use a generic GPU kernel as fallback when running on GPU.
WARNING:tensorflow:Layer lstm_1 will not use cuDNN kernels since it doesn't meet the criteria. It will use a generic GPU kernel as fallback when running on GPU.

Out[181]...

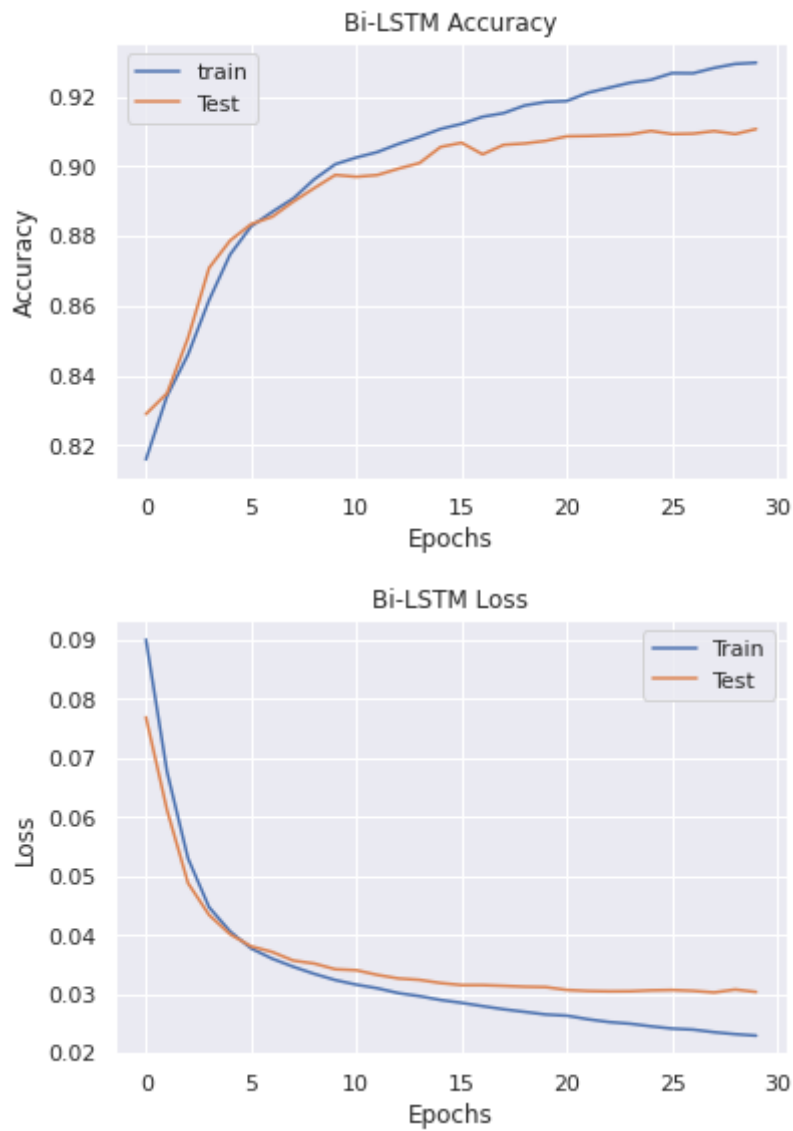


In [40]:

```
history = model.fit(train_vec, train_ent_vec,
                    batch_size=batch_size,
                    epochs=epochs,
                    validation_data=[test_vec, test_ent_vec],
                    verbose=0, callbacks=[TqdmCallback(verbose=1)])
```

In [43]:

```
plot_acc(history)
plot_loss(history)
```



In [44]:

```
y_true, y_pred = [], []
test_pred = decode(model.predict(test_vec))
for i in range(len(test_entities)):
    y_true+=test_entities[i]
    y_pred+=test_pred[i][:len(test_entities[i])]
print(classification_report(y_true, y_pred))
```

210/210 [=====] - 24s 110ms/step

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

LOC	0.67	0.46	0.55	4019
MISC	0.51	0.43	0.47	2186
O	0.95	0.97	0.96	81082
ORG	0.61	0.54	0.58	4588
PER	0.76	0.82	0.79	5922
accuracy			0.91	97797
macro avg	0.70	0.65	0.67	97797
weighted avg	0.90	0.91	0.91	97797

```
In [45]: idx = 3
_x_wrd, _x_lab, _x_vec = test_sent[idx], test_entities[idx], test_vec[idx]
_y_pred = decode(model.predict(np.array([_x_vec,])))[0][:len(_x_wrd)]
org = [(_x_wrd[i], _x_lab[i]) for i in range(len(_x_lab))]
item = [(_x_wrd[i], _y_pred[i]) for i in range(len(_x_wrd))]
```

1/1 [=====] - 0s 198ms/step

```
In [46]: doc = displacy_doc(org, "Original NER")
html = displacy.render(doc, style="ent", manual=True, jupyter=True)
```

Original NER

Their o stay o on o top o , o though o , o may o be o
short-lived o as o title o rivals o Essex **ORG** , o Derbyshire **ORG**
and o Surrey **ORG** all o closed o in o on o victory o while o
Kent **ORG** made o up o for o lost o time o in o their o
rain-affected o match o against o Nottinghamshire **ORG** . o

```
In [47]: doc = displacy_doc(item, "Generated NER")
html = displacy.render(doc, style="ent", manual=True, jupyter=True)
```

Generated NER

Their o stay o on o top o , o though o , o may o be o
short-lived o as o title o rivals o Essex **ORG** , o Derbyshire **ORG**
and o Surrey **ORG** all o closed o in o on o victory o while o
Kent **PER** made o up o for o lost o time o in o their o
rain-affected o match o against o Nottinghamshire **LOC** . o

LSTM Model on Synthetic Dataset

```
In [99]: maxlen=25
synth_train_sent = [[x[0] for x in sent] for sent in synth_train_data]
synth_test_sent = [[x[0] for x in sent] for sent in synth_test_data]
train_sent_vec = pad_sequences(word2vec(synth_train_sent), maxlen=maxlen, padding=
test_sent_vec = pad_sequences(word2vec(synth_test_sent), maxlen=maxlen, padding=
train_entities = [[x[1] for x in sent] for sent in synth_train_data]
test_entities = [[x[1] for x in sent] for sent in synth_test_data]
train_ent_vec = pad_sequences(encode(train_entities), maxlen=maxlen, padding='post
value= [0, 0, 0, 0, 1])
test_ent_vec = pad_sequences(encode(test_entities), maxlen=maxlen, padding='post
value= [0, 0, 0, 0, 1])
```

```
<ipython-input-31-74d0e4c39a73>:8: VisibleDeprecationWarning: Creating an ndarray from ragged nested sequences (which is a list-or-tuple of lists-or-tuples-or ndarrays with different lengths or shapes) is deprecated. If you meant to do this, you must specify 'dtype=object' when creating the ndarray.
return np.array(out)
```

In [125...

```
if GLOVE:
    model = get_bilstm_lstm_model()
    train_vec = train_sent_vec
    test_vec = test_sent_vec
else:
    pass
```


```
WARNING:tensorflow:Layer lstm_30 will not use cuDNN kernels since it doesn't meet the criteria. It will use a generic GPU kernel as fallback when running on GPU.
WARNING:tensorflow:Layer lstm_30 will not use cuDNN kernels since it doesn't meet the criteria. It will use a generic GPU kernel as fallback when running on GPU.
WARNING:tensorflow:Layer lstm_30 will not use cuDNN kernels since it doesn't meet the criteria. It will use a generic GPU kernel as fallback when running on GPU.
WARNING:tensorflow:Layer lstm_31 will not use cuDNN kernels since it doesn't meet the criteria. It will use a generic GPU kernel as fallback when running on GPU.
```

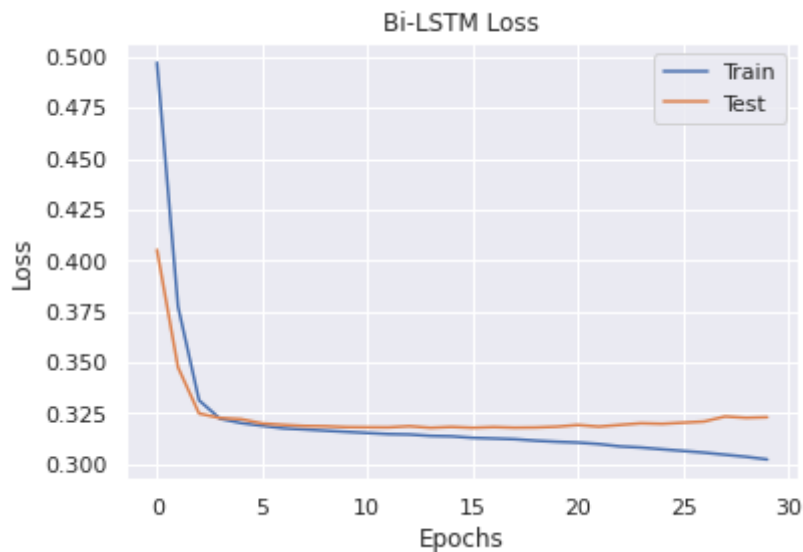
In [126...

```
history = model.fit(train_vec, train_ent_vec,
                    batch_size=batch_size,
                    epochs=epochs,
                    validation_data=[test_vec, test_ent_vec],
                    verbose=0, callbacks=[TqdmCallback(verbose=1)])
```

In [127...

```
plot_acc(history)
plot_loss(history)
```





In [128...

```
y_true, y_pred = [], []
test_pred = decode(model.predict(test_vec))
for i in range(len(test_entities)):
    y_true+=test_entities[i]
    y_pred+=test_pred[i][:len(test_entities[i])]
print(classification_report(y_true, y_pred))
```

210/210 [=====] - 6s 26ms/step

/usr/local/lib/python3.8/dist-packages/sklearn/metrics/_classification.py:1318: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.

```
_warn_prf(average, modifier, msg_start, len(result))
```

	precision	recall	f1-score	support
LOC	0.25	0.01	0.01	4516
MISC	0.00	0.00	0.00	2549
O	0.86	0.99	0.92	95958
ORG	0.31	0.20	0.24	5287
PER	0.31	0.07	0.12	6004
accuracy			0.84	114314
macro avg	0.35	0.25	0.26	114314
weighted avg	0.77	0.84	0.79	114314

/usr/local/lib/python3.8/dist-packages/sklearn/metrics/_classification.py:1318: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.

```
_warn_prf(average, modifier, msg_start, len(result))
```

/usr/local/lib/python3.8/dist-packages/sklearn/metrics/_classification.py:1318: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.

```
_warn_prf(average, modifier, msg_start, len(result))
```

In [159...

```
idx = 3
_x_wrd, _x_lab, _x_vec = synth_test_sent[idx], test_entities[idx], test_vec[idx]
_y_pred = decode(model.predict(np.array([_x_vec,])))[0][:len(_x_wrd)]
```

```
org = [(_x_wrd[i], _x_lab[i]) for i in range(len(_x_lab))]  
item = [(_x_wrd[i], _y_pred[i]) for i in range(len(_x_wrd))]
```

1/1 [=====] - 0s 71ms/step

In [160...

```
doc = displacy_doc(org, "Generated NER")  
html = displacy.render(doc, style="ent", manual=True, jupyter=True)
```

Generated NER

finally o the o Two o of o , o (o Friday o AND o is o
who o the o Jerry PER Bevan PER David PER 0.38 o Hungary
LOC of o the o surprise o

In [161...

```
doc = displacy_doc(item, "Predicted NER")  
html = displacy.render(doc, style="ent", manual=True, jupyter=True)
```

Predicted NER

finally o the o Two o of o , o (o Friday o AND o is o
who o the o Jerry o Bevan ORG David ORG 0.38 ORG Hungary
ORG of o the o surprise o